Free Questions for AIF-C01

Shared by Langley on 02-09-2025

For More Free Questions and Preparation Resources

Check the Links on Last Page

Question 1

Question Type: MultipleChoice

A hospital developed an AI system to provide personalized treatment recommendations for patients. The AI system must provide the rationale behind the recommendations and make the insights accessible to doctors and patients.

Which human-centered design principle does this scenario present?

Options:

A- Explainability
B- Privacy and security
C- Fairness
D- Data governance

Answer:

A

Explanation:

Explainability refers to the ability of an AI system to make its decision-making process clear and understandable to humans.

A is correct:

'Explainability is crucial for human-centered AI, especially in healthcare, to ensure that doctors and patients understand the rationale behind AI-driven recommendations.' (Reference: AWS Responsible AI)

'Explainability is crucial for human-centered AI, especially in healthcare, to ensure that doctors and patients understand the rationale behind AI-driven recommendations.' (Reference: AWS Responsible AI)

B relates to protecting data, not explanations.

C is about treating groups equally.

D is about managing data lifecycle, not providing rationales.

Question 2

Question Type: MultipleChoice

A financial company is using ML to help with some of the company's tasks.

Which option is a use of generative AI models?

Options:

A- Summarizing customer complaints
B- Classifying customers based on product usage
C- Segmenting customers based on type of investments
D- Forecasting revenue for certain products

Answer:

A

Explanation:

Generative AI models (such as large language models) are designed to generate new content, such as text, summaries, images, and more. Summarizing text---like customer complaints---is a classic application of generative AI.

A is correct:

'Text summarization is a core generative AI use case, as it involves generating new, concise content from a larger body of text.' (Reference: AWS Generative AI Use Cases)

'Text summarization is a core generative AI use case, as it involves generating new, concise content from a larger body of text.' (Reference: AWS Generative AI Use Cases)

B and C are standard ML classification/segmentation tasks.

D is a regression/prediction task, not generative.

Question 3

Question Type: Hotspot

A company wants more customized responses to its generative AI models' prompts.

Select the correct customization methodology from the following list for each use case. Each use case should be selected one time. (Select THREE.)

* Continued pre-training

* Data augmentation

* Model fine-tuning

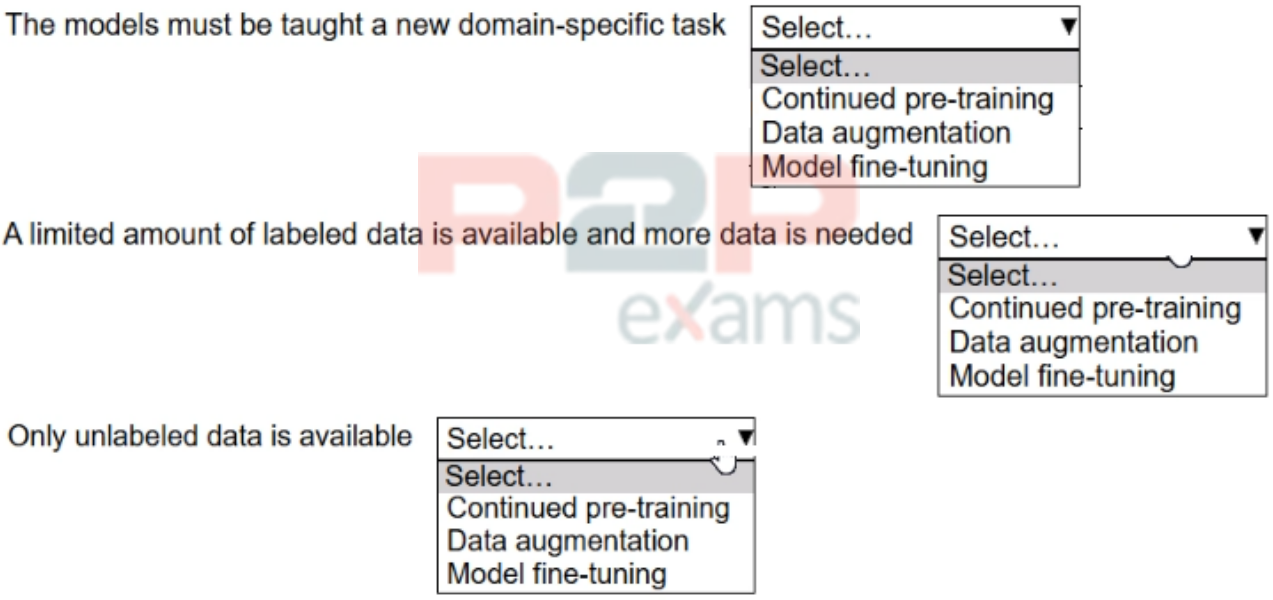

Answer:

See the Answer in the Premium Version!

Question 4

Question Type: MultipleChoice

A financial company is developing a fraud detection system that flags potential fraud cases in credit card transactions. Employees will evaluate the flagged fraud cases. The company wants to minimize the amount of time the employees spend reviewing flagged fraud cases that are not actually fraudulent.

Which evaluation metric meets these requirements?

Options:

A- Recall

B- Accuracy

C- Precision
D- Lift chart

Answer:

C

Explanation:

Precision is the metric that measures the proportion of true positives (actual frauds) among all flagged positives (flagged frauds). High precision ensures that most of the flagged cases are truly fraudulent, minimizing the number of false positives employees must review.

C is correct:

'Precision is the ratio of true positives to all predicted positives, and it answers: 'Of all the cases flagged as fraud, how many were actually fraud?' High precision means fewer non-fraudulent cases are sent for manual review.' (Reference: AWS ML Concepts -- Precision and Recall, AWS Certified AI Practitioner Study Guide)

'Precision is the ratio of true positives to all predicted positives, and it answers: 'Of all the cases flagged as fraud, how many were actually fraud?' High precision means fewer non-fraudulent cases are sent for manual review.' (Reference: AWS ML Concepts -- Precision and Recall, AWS Certified AI Practitioner Study Guide)

A (Recall) measures how many actual frauds are caught, but does not minimize false positives.

B (Accuracy) can be misleading in imbalanced datasets (like fraud detection).

Question 5

Question Type: MultipleChoice

A company that uses multiple ML models wants to identify changes in original model quality so that the company can resolve any issues.

Which AWS service or feature meets these requirements?

Options:

A- Amazon SageMaker JumpStart
B- Amazon SageMaker HyperPod
C- Amazon SageMaker Data Wrangler

D- Amazon SageMaker Model Monitor

Answer:

D

Explanation:

Amazon SageMaker Model Monitor is specifically designed to automatically detect and alert on changes in model quality, such as data drift, prediction drift, or other anomalies in model performance once deployed.

D is correct:

'Amazon SageMaker Model Monitor continuously monitors the quality of machine learning models in production. It automatically detects concept drift, data drift, and other quality issues, enabling teams to take corrective actions.' (Reference: Amazon SageMaker Model Monitor Documentation, AWS Certified AI Practitioner Study Guide)

'Amazon SageMaker Model Monitor continuously monitors the quality of machine learning models in production. It automatically detects concept drift, data drift, and other quality issues, enabling teams to take corrective actions.' (Reference: Amazon SageMaker Model Monitor Documentation, AWS Certified AI Practitioner Study Guide)

A (JumpStart) provides prebuilt solutions and models, not monitoring.

B (HyperPod) is for large-scale training, not model monitoring.

C (Data Wrangler) is for data preparation, not ongoing model quality monitoring.

Question 6

Question Type: Hotspot

Select the correct AI term from the following list for each statement. Each AI term should be selected one time. (Select THREE.)

* AI

* Deep learning

* ML

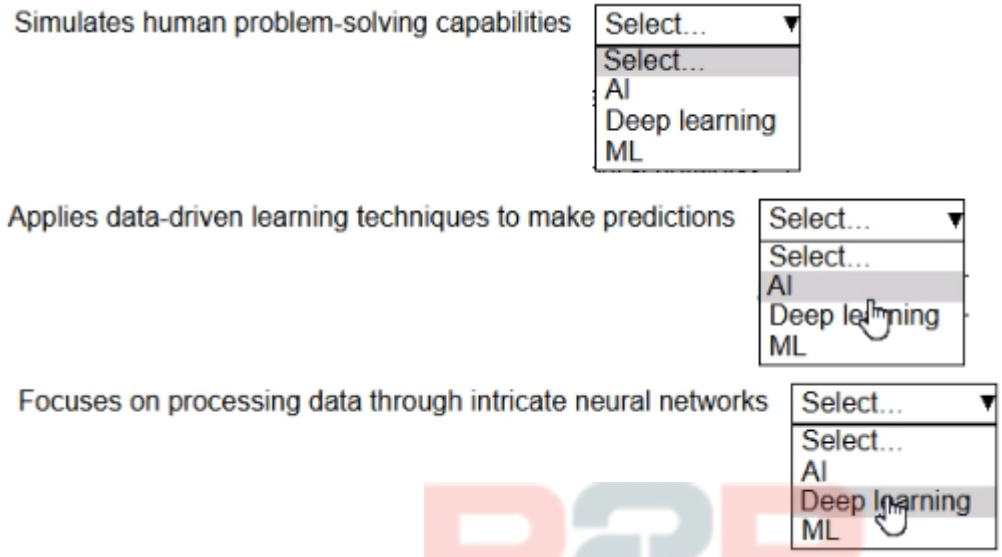

Answer:

See the Answer in the Premium Version!

Question 7

Question Type: MultipleChoice

A hospital is developing an AI system to assist doctors in diagnosing diseases based on patient records and medical images. To comply with regulations, the sensitive patient data must not leave the country the data is located in.

Which data governance strategy will ensure compliance and protect patient privacy?

Options:

A- Data residency
B- Data quality
C- Data discoverability
D- Data enrichment

Answer:

A

Explanation:

Data residency is the principle and practice of ensuring that data remains within a specific geographic location or jurisdiction, often to comply with local regulations and privacy laws (such as HIPAA, GDPR, or national healthcare laws). Data residency policies prevent sensitive data (such as patient records) from being transferred or accessed outside the designated country, thus protecting privacy and ensuring regulatory compliance.

A is correct:

''Data residency refers to where data is stored geographically, and often organizations need to ensure that certain data does not leave a particular country or region to comply with legal or regulatory requirements.'' (Reference: AWS Data Residency Whitepaper, AWS Responsible AI & Data Privacy)

''Data residency refers to where data is stored geographically, and often organizations need to ensure that certain data does not leave a particular country or region to comply with legal or regulatory requirements.'' (Reference: AWS Data Residency Whitepaper, AWS Responsible AI & Data Privacy)

B (data quality) refers to the accuracy and reliability of data, not its location.

C (data discoverability) is about being able to find and access data, not restricting its movement.

D (data enrichment) is about enhancing data with additional information.

''Maintaining data residency is critical in healthcare and regulated industries to ensure data does not leave the prescribed jurisdiction.'' (Reference: AWS Data Residency)

''Maintaining data residency is critical in healthcare and regulated industries to ensure data does not leave the prescribed jurisdiction.'' (Reference: AWS Data Residency)