Free Questions for MLA-C01

Shared by Donaldson on 02-09-2025

For More Free Questions and Preparation Resources

Check the Links on Last Page

Question 1

Question Type: MultipleChoice

An ML engineer needs to deploy ML models to get inferences from large datasets in an asynchronous manner. The ML engineer also needs to implement scheduled monitoring of the data quality of the models. The ML engineer must receive alerts when changes in data quality occur.

Which solution will meet these requirements?

Options:

A- Deploy the models by using scheduled AWS Glue jobs. Use Amazon CloudWatch alarms to monitor the data quality and to send alerts.
B- Deploy the models by using scheduled AWS Batch jobs. Use AWS CloudTrail to monitor the data quality and to send alerts.
C- Deploy the models by using Amazon Elastic Container Service (Amazon ECS) on AWS Fargate. Use Amazon EventBridge to monitor the data quality and to send alerts.
D- Deploy the models by using Amazon SageMaker batch transform. Use SageMaker Model Monitor to monitor the data quality and to send alerts.

Answer:

D

Question 2

Question Type: MultipleChoice

A company wants to develop an ML model by using tabular data from its customers. The data contains meaningful ordered features with sensitive information that should not be discarded. An ML engineer must ensure that the sensitive data is masked before another team starts to build the model.

Which solution will meet these requirements?

Options:

A- Use Amazon Made to categorize the sensitive data.
B- Prepare the data by using AWS Glue DataBrew.
C- Run an AWS Batch job to change the sensitive data to random values.

D- Run an Amazon EMR job to change the sensitive data to random values.

Answer:

B

Question 3

Question Type: MultipleChoice

A company is using Amazon SageMaker and millions of files to train an ML model. Each file is several megabytes in size. The files are stored in an Amazon S3 bucket. The company needs to improve training performance.

Which solution will meet these requirements in the LEAST amount of time?

Options:

A- Transfer the data to a new S3 bucket that provides S3 Express One Zone storage. Adjust the training job to use the new S3 bucket.
B- Create an Amazon FSx for Lustre file system. Link the file system to the existing S3 bucket. Adjust the training job to read from the file system.
C- Create an Amazon Elastic File System (Amazon EFS) file system. Transfer the existing data to the file system. Adjust the training job to read from the file system.
D- Create an Amazon ElastiCache (Redis OSS) cluster. Link the Redis OSS cluster to the existing S3 bucket. Stream the data from the Redis OSS cluster directly to the training job.

Answer:

B

Question 4

Question Type: MultipleChoice

A company is running ML models on premises by using custom Python scripts and proprietary datasets. The company is using PyTorch. The model building requires unique domain knowledge. The company needs to move the models to AWS.

Which solution will meet these requirements with the LEAST effort?

Options:

A- Use SageMaker built-in algorithms to train the proprietary datasets.
B- Use SageMaker script mode and premade images for ML frameworks.
C- Build a container on AWS that includes custom packages and a choice of ML frameworks.
D- Purchase similar production models through AWS Marketplace.

Answer:

B

Question 5

Question Type: MultipleChoice

An ML engineer is using Amazon SageMaker to train a deep learning model that requires distributed training. After some training attempts, the ML engineer observes that the instances are not performing as expected. The ML engineer identifies communication overhead between the training instances.

What should the ML engineer do to MINIMIZE the communication overhead between the instances?

Options:

A- Place the instances in the same VPC subnet. Store the data in a different AWS Region from where the instances are deployed.
B- Place the instances in the same VPC subnet but in different Availability Zones. Store the data in a different AWS Region from where the instances are deployed.
C- Place the instances in the same VPC subnet. Store the data in the same AWS Region and Availability Zone where the instances are deployed.
D- Place the instances in the same VPC subnet. Store the data in the same AWS Region but in a different Availability Zone from where the instances are deployed.

Answer:

C

Question 6

Question Type: MultipleChoice

A company's ML engineer has deployed an ML model for sentiment analysis to an Amazon SageMaker endpoint. The ML engineer needs to explain to company stakeholders how the model makes predictions.

Which solution will provide an explanation for the model's predictions?

Options:

A- Use SageMaker Model Monitor on the deployed model.
B- Use SageMaker Clarify on the deployed model.
C- Show the distribution of inferences from A/ testing in Amazon CloudWatch.
D- Add a shadow endpoint. Analyze prediction differences on samples.

Answer:

B